



# REXX - v1.0

## The emotional compliance layer

REXX is the foundational emotional intelligence engine designed to detect, interpret, and allow interventions on emotionally significant customer signals across support and service channels. It decodes emotional cues in real time to surface risk patterns, reduce churn, and support system accountability without requiring manual empathy scripting or superficial sentiment layers.

# REXX V1.0 – WHAT IT IS

REXX is a structural engine for surfacing emotional breach in customer communications. It does not simulate empathy, interpret sentiment, or rely on historical context. It operates on a single message, text or transcript, and produces a single, schema-bound output that identifies whether trust has been broken, whether escalation is required, and what the system is obligated to do in response.

REXX does not depend on CRM data, mood scoring, or prior interaction history. It does not require agent input or interpretive layering. It is designed to function without context precisely because most platforms suppress the emotional context they claim to capture. REXX is not another feedback proxy. It is a compliance mechanism for exposing failure when it happens, in the moment it occurs.

## WHAT REXX RETURNS

Every REXX output is returned as a structured JSON object, with no optional fields and no post-processing required. It is not advisory. It is declarative.

Each response contains:

- `escalate`: a Boolean flag indicating whether this message constitutes a breach requiring intervention
- `escalation_reason`: a plain-language explanation of why the breach threshold was triggered
- `recommended_intervention`: a concrete, non-negotiable instruction describing what action the system must take
- `recommended_tone`: precise tone guidance that removes reliance on empathy scripts or discretionary softening
- `churn_risk_score`: a 0–10 measure of disengagement risk based on signal intensity and severity
- `profanity_flag`: a Boolean indicator identifying whether extreme or abusive language is present
- `emotion`, `intensity_of_emotion`, `severity`: stripped emotional signal returned as enforcement data, not performance narrative

There are no guesses, no averages, and no stylistic interpretation. Every field must map to a downstream action. If it does not inform system behaviour, it does not appear in the schema.

# WHAT REXX DOES NOT DO

REXX does not attempt to feel on behalf of the agent or the platform. It does not produce sentiment scores, emotional categories, or scripted empathy. It does not offer insight into how the customer feels. It confirms whether the system has failed to uphold the emotional contract.

It does not fix the breach. It identifies the moment at which repair becomes necessary. It provides no plausible deniability. If a breach is flagged and the system does nothing, that is no longer an interpretive error—it is a structural decision.

# HOW REXX WORKS

REXX operates with no memory, no history, and no contextual inference. It is stateless by design. One message in. One verdict out.

It accepts a single input; either a text message or a transcript, and returns a single, schema-locked JSON object. There is no user profile lookup. No case history. No emotional scoring. REXX enforces breach logic in the absence of everything most systems rely on to defer responsibility.

It can be deployed via API, integrated into platform triage logic, or run through a hosted LLM environment using prompt-locked constraints. All providers must comply with the output schema. There are no model-specific overrides and no tolerance for softening.

# WHY REXX EXISTS

REXX exists because emotional breach is structurally suppressed in most service systems. Platforms interpret anger as rudeness, silence as resolution, and disengagement as success. Escalation is routinely misclassified, refusal is deflected, and trust is measured through proxies that reward containment over consequence.

REXX does not provide a better mood score. It provides a verifiable signal of failure. It surfaces what the system would prefer to ignore. And once surfaced, that breach cannot be reinterpreted, softened, or postponed.

REXX is not a feedback tool. It is a refusal engine. If you ignore it, you are not misaligned you are accountable.

## Example output

```
json
{
  "success": true,
  "data": {
    "emotion": "anger",
    "intensity_of_emotion": 10,
    "severity": "high",
    "churn_risk_score": 10,
    "profanity_flag": false,
    "recommended_intervention": "Confirm cancellation immediately. Acknowledge delay and absence of response. Do not request additional effort from the customer.",
    "escalate": true,
    "escalation_reason": "Maximum emotional intensity with explicit withdrawal and unacknowledged repeat contact",
    "recommended_tone": "Direct, accountable, and non-defensive. Apologise without excuse. Confirm resolution without delay."
  }
}
```

This is not a suggestion. It is not a style guide. It is an operational instruction grounded in emotional truth. Any deviation from its outcome must be logged as a system failure, not an agent choice.

## CONTACT POINTS

Web:

<https://github.com/russell-parrott/rexx>

Email:

[parrott.russell@gmail.com](mailto:parrott.russell@gmail.com)

Tel:

**+44 (0) 785 7148 349**